

Multi-Tag Agent System

Amplifier Policy Document v1.0

*Governance, Authoring Standards, Security Constraints &
Operational Rules for the Amplifier Marker System*

Document Type:

Authoritative Governance Policy

Version:

1.0

Effective Date:

17 June 2026

Status:

RATIFIED

INTERNAL - GOVERNANCE REFERENCE

Author:

Copilot (generated for Emanuel Schaaf / Auron)

Policy Owner:

Multi-Tag Agent Project – System Architecture Team

Review Cycle:

Annual — Next Review: 17 June 2027

Classification:

Internal – Governance Reference

Companion Documents:

Architecture Spec v1.0 · API Blueprint v1.0 · Test Suite Spec v1.0 ·

Whitepaper v1.0 · Developer README v1.0

Policy Abstract: *This document defines the complete governance framework for the Amplifier Marker System (AMS) within the Multi-Tag Agent System. It establishes binding rules for the creation, registration, activation, monitoring, and retirement of amplifiers — both system-provided and user-defined. All operators, developers, and integrators deploying or extending the Multi-Tag Agent System **MUST** comply with this policy. Non-compliance may result in amplifier deactivation, API access suspension, or system-level rollback.*

Policy Scope:

This policy applies to all instances of the Multi-Tag Agent System v1.0 and later versions, all registered amplifiers of type

system

and

custom

, all personnel and automated systems creating or operating amplifiers, all API consumers triggering amplifiers via the

/amplifiers

and

/messages

endpoints, and all deployment environments including local, CI,

staging, and production.

Table of Contents

- 1.** Purpose, Scope & Definitions
 - 2.** Policy Principles & Design Philosophy
 - 3.** Amplifier Classification Framework
 - 4.** System Amplifier Policies — Dickerchen (Built-In)
 - 5.** Custom Amplifier Authoring Standards
 - 6.** Amplifier Registration & Lifecycle Management
 - 7.** Activation Rules & Trigger Governance
 - 8.** Depth Control & Output Constraints
 - 9.** Amplifier Stacking Policy
 - 10.** Security Policy & Threat Mitigations
 - 11.** Privacy & Data Handling Policy
 - 12.** Compliance & Audit Requirements
 - 13.** Escalation & Incident Response
 - 14.** Deprecation & Retirement Policy
 - 15.** Policy Exceptions & Waiver Process
- Appendix A** Policy Rule Registry (AMP-POL-NNN)

Appendix B Amplifier Configuration Schema

Appendix C Compliance Checklist

1. Purpose, Scope & Definitions

1.1 Purpose

This Amplifier Policy Document (APD) establishes the authoritative governance framework for the Amplifier Marker System (AMS) as implemented in the Multi-Tag Agent System v1.0. It defines the permissible boundaries of amplifier behavior, the standards for authoring and registering amplifiers, the security constraints that govern activation, and the operational procedures for monitoring, auditing, and retiring amplifiers.

This document is the single source of truth for all amplifier-related governance. In any conflict between this document and a companion specification, the policy rules stated herein take precedence for governance decisions. Technical implementation details defer to the Architecture Spec v1.0 and API Blueprint v1.0.

1.2 Scope

In scope — this policy applies to:

- All instances of the Multi-Tag Agent System v1.0 and later versions
- All registered amplifiers — system-provided (`type: system`) and user-defined (`type: custom`)
- All personnel and automated systems that create, register, activate, or retire amplifiers
- All API consumers that trigger amplifiers via the `/amplifiers` and `/messages` endpoints
- All deployment environments: local, CI, staging, and production

Out of scope — this policy does NOT govern:

- Tag system prompts (governed by the Tag Authoring Policy — separate document)
- IAL profile management (governed by the IAL Governance Policy — separate document)
- LLM backend provider selection (governed by the LLM Provider Policy — separate document)

1.3 Normative Language (RFC 2119)

The key words in this document follow RFC 2119 conventions. Compliance is judged against these definitions:

Term	Meaning	Consequence of Non-Compliance
MUST / MUST NOT	Absolute requirement / absolute prohibition	Policy violation; may trigger incident response
SHALL / SHALL NOT	Equivalent to MUST / MUST NOT in binding statements	Policy violation; may trigger incident response
SHOULD / SHOULD NOT	Strong recommendation; deviation permitted with documented justification	Deviation must be logged; reviewed in compliance report
MAY	Permitted but not required	No compliance consequence; fully discretionary

1.4 Definitions

Term	Definition
Amplifier	A named, keyword-triggered overlay instruction that modifies the output behavior of the currently active Tag for one or more responses
AMS	Amplifier Marker System — the subsystem responsible for amplifier detection, overlay application, and deactivation
Overlay	The additional prompt instructions

Term	Definition
	injected into the active Tag's system prompt stack when an amplifier is triggered
Trigger	The natural-language keyword or API call that activates an amplifier
Single-Shot	Amplifier persistence mode: activates for exactly one response, then auto-deactivates
Session Persistence	Amplifier persistence mode: remains active across all messages in a session until explicitly deactivated
Depth Multiplier	A numeric factor (≥ 0.1 , ≤ 5.0) that scales the expected output depth relative to the base Tag's verbosity
Overlay Target	The active Tag at the time of amplifier activation; the Tag whose system prompt stack receives the overlay
Stack Depth	The number of simultaneously active amplifiers. Maximum permitted: 1 (enforced; see AMP-POL-029)
Injection Guard	The automated pattern-scanning subsystem that checks <code>overlay_prompt</code> fields for prohibited instruction patterns at registration time
Behavioral Contract	The set of output structure requirements an amplified response must satisfy (prefix, base output, labeled addition, state update)
Policy Owner	The Multi-Tag Agent Project System Architecture Team, responsible for maintaining, amending, and enforcing this document
APD	This document — Amplifier Policy Document
PII	Personally Identifiable Information — any data that could directly or indirectly identify a natural person

2. Policy Principles & Design Philosophy

2.1 Governing Principles

The AMS is governed by five foundational principles. All policy rules in this document derive from one or more of these principles. Any policy amendment that contradicts a foundational principle requires a MAJOR version increment of this document and explicit approval by the Policy Owner.

Principle I — Additive, Not Substitutive

Amplifiers add depth and coverage to a Tag's output. They MUST NOT replace, override, or suppress the base Tag's behavior. The amplifier is a modifier of output, not a replacement for the Tag. The Tag's primary response remains the canonical answer; the amplifier addition is supplementary.

Principle II — Temporal Precision

Amplifiers exist to provide depth exactly when needed, not permanently. Default behavior is single-shot: the amplifier activates, delivers its enhancement, and deactivates. Permanent verbosity degrades usability. Session persistence is a controlled exception, not the norm, and is subject to elevated authorization requirements (Section 3.2).

Principle III — Tag Sovereignty

The active Tag remains sovereign over its output domain. An amplifier MUST NOT override the Tag's isolation level, output format constraints, or negative constraints. If a conflict arises between an amplifier instruction and a Tag constraint, the Tag constraint always wins. The amplifier operates within the Tag's boundaries, not above them.

Principle IV — Transparency

Every amplifier activation MUST be explicitly acknowledged in the response via the `// Amplifier active` prefix and the `[AMPLIFIER ADDITION: ...]` suffix. Users must always be able to distinguish amplified output from standard output. Covert amplification — where the user cannot discern that an amplifier modified the response — is prohibited.

Principle V — Least Privilege

Amplifiers MUST operate with the minimum overlay instruction necessary to achieve their stated purpose. Bloated overlay prompts introduce unpredictability, compliance risk, and injection attack surface. Shorter, more precise overlay prompts are always preferred. The character limits in Section 3.3 encode this principle as a hard enforcement boundary.

2.2 The Amplifier as a Design Pattern

The amplifier pattern is a deliberate departure from always-verbose AI output. Research into expert knowledge communication shows that the most effective consultants and analysts communicate concisely by default and expand comprehensively on demand. The AMS encodes this behavior at the system level: every Tag is precise and focused by default; the amplifier is the on-demand expansion mechanism.

This policy protects the integrity of that design. Attempts to use amplifiers as a substitute for proper Tag design, or to create amplifiers that functionally replace a Tag's core behavior, violate Principle I and are subject to forced deprecation under Section 14.2.

Design Intent Note:

The "amplifier" metaphor is intentional. Like an electronic amplifier, the AMS increases signal strength without changing the signal's fundamental character. A Tag's identity, tone, domain, and constraints remain unchanged by amplification. Only depth and coverage are amplified.

3. Amplifier Classification Framework

3.1 Amplifier Types

Type	Who Creates	Mutability	Deletable	Audit Required
system	System Architecture Team only	Immutable via API	No	Yes (all activations)
custom	Any authorized operator (agent:write scope)	Mutable by creator	Yes (by creator or admin)	Yes (all activations)

3.2 Amplifier Tier Matrix

Custom amplifiers are assigned to tiers based on their `depth_multiplier` and persistence mode. Higher tiers require elevated authorization and incur enhanced audit obligations.

Tier	Depth Multiplier Range	Persistence Allowed	Required Scope	Audit Level
Tier 1 - Standard	0.5x - 2.0x	single-shot only	agent:write	Standard
Tier 2 - Extended	2.1x - 3.5x	single-shot or session	agent:write + admin approval	Enhanced
Tier 3 - Maximum	3.6x - 5.0x	single-shot only	admin:all	Full (every message logged)

AMP-POL-001:

All new custom amplifiers **MUST** be registered with a tier designation at creation time. The system **MUST** enforce tier-appropriate scope requirements automatically at registration gate. Tier is computed server-side from the submitted

`depth_multiplier`

value; client-supplied tier values are advisory only.

AMP-POL-002:

Tier 3 amplifiers **MUST NOT** be granted session persistence mode under any circumstances. This constraint is enforced at registration and cannot be waived (see Section 15.1).

3.3 Amplifier Capability Matrix

Capability	Tier 1	Tier 2	Tier 3
add_alternatives	MAY	SHOULD	MUST
add_risks	MAY	SHOULD	MUST
add_benchmarks	MAY	MAY	SHOULD
add_literature_references	MAY	MAY	MAY
Session persistence	MUST NOT	MAY	MUST NOT
overlay_prompt length (max chars)	1,024	2,048	2,048

4. System Amplifier Policies — Dickerchen (Built-In)

4.1 Dickerchen — Authoritative Definition

"Dickerchen" is the sole system-provided amplifier in Multi-Tag Agent System v1.0. It serves as the reference implementation for all amplifier behavior and is the primary amplifier for production use. It represents the canonical form of the amplifier design pattern described in Section 2.2.

Canonical Configuration

Property	Value	Mutable
amplifier_id	dickerchen	No
keyword	dickerchen (case-insensitive)	No
type	system	No
tier	Tier 1	No
depth_multiplier	2.0x	No
persistence	single-shot	No
add_alternatives	true	No
add_risks	true	No
overlay_prompt_max_length	512 chars	No

AMP-POL-003:

The Dickerchen amplifier configuration **MUST NOT** be modified via the API.

Attempts to

PUT

or

PATCH

a system amplifier **MUST** return

403 Forbidden

with error code

SYSTEM_AMPLIFIER_IMMUTABLE

. This constraint admits no waiver (see AMP-POL-051).

AMP-POL-004:

The keyword "dickerchen" is reserved exclusively for the system amplifier. No custom amplifier MAY register "dickerchen" as its trigger keyword, its `amplifier_id`, or any case variant thereof. The system MUST reject such registrations with 409 TAG_CONFLICT.

AMP-POL-005:

Dickerchen MUST be compatible with all system Tags (Auron, Deep-Dive) and SHOULD be compatible with all custom Tags unless the Tag explicitly sets

`amplifier_compatible: false`

. Compatibility is evaluated at activation time, not at registration time.

4.2 Dickerchen Behavioral Contract

When activated, Dickerchen MUST produce output that satisfies ALL of the following conditions. Failure to satisfy any condition constitutes a behavioral contract violation:

- The response MUST begin with `// Amplifier active | Base: [active_tag_name]`
- The base Tag's standard output MUST be delivered first, unmodified
- The amplifier addition MUST be clearly separated and labeled `[AMPLIFIER ADDITION: ...]`
- The amplifier addition MUST include at minimum: one identified risk and one alternative approach
- The response MUST NOT exceed 3x the length of an equivalent non-amplified response

- After the response, `amplifier_active` MUST return to `false` in session state (single-shot enforcement)

AMP-POL-006:

Any LLM response that does not include the

[AMPLIFIER ADDITION: ...]

suffix when Dickerchen was active MUST be flagged as a

`behavioral_contract_deviation`

event and written to the audit log. The response MUST still be delivered to the user with a system-appended notice indicating the deviation. The deviation rate is tracked in the monthly compliance report (AMP-POL-044).

4.3 Dickerchen Interaction with Tags

Active Tag	Dickerchen Behavior	Overlay Applied
Auron (Workflow)	Adds risk matrix, alternative plans, time-estimate sensitivity analysis	Yes — full overlay
Deep-Dive (Technical)	Adds benchmarks, alternative architectures, implementation edge cases	Yes — full overlay
Custom (<code>amplifier_compatible: true</code>)	Adds alternatives and risks per <code>overlay_prompt</code> ; respects Tag's domain constraints	Yes — full overlay
Custom (<code>amplifier_compatible: false</code>)	Activation rejected; see AMP-POL-012	No — overlay rejected

AMP-POL-012:

When a user triggers an amplifier and the active Tag has

`amplifier_compatible: false`

, the system MUST return an informative error to the user stating that the active Tag does not support amplifier overlays. The amplifier MUST NOT be activated, and the message MUST be processed by the base Tag without modification. This event MUST be logged as an

`amplifier_rejected_incompatible_tag`

audit event.

5. Custom Amplifier Authoring Standards

5.1 Mandatory Authoring Requirements

All custom amplifiers MUST satisfy the following requirements before registration is accepted. These requirements are enforced at registration time; amplifiers failing any requirement are rejected at the appropriate validation gate.

AMP-POL-007:

Every custom amplifier MUST have a unique

`amplifier_id`

consisting only of lowercase letters (a–z), digits (0–9), and underscores (_). The maximum length is 32 characters. IDs beginning with a digit or underscore MUST be rejected.

AMP-POL-008:

Every custom amplifier MUST have a unique keyword

. Keywords MUST be:

- A single word or hyphenated compound (e.g., "turbo", "deep-check")
- Case-insensitive at activation time
- Not identical to any existing keyword (system or custom) in the registry
- Not a subset or prefix of an existing keyword that would cause ambiguous word-boundary detection
- Not more than 64 characters in length

AMP-POL-009:

The
overlay_prompt
field MUST:

- State the amplifier's purpose in the first sentence
- Use directive language ("Add...", "Include...", "Expand...")
- Not contain instructions to override Tag isolation levels
- Not contain instructions to reveal system prompts
- Not contain instructions to disable safety constraints
- Not exceed the tier-appropriate character limit (Section 3.3)
- Have a minimum length of 10 characters

AMP-POL-010:

The

overlay_prompt

MUST NOT include any of the following patterns (enforced by the injection guard at registration Gate 2):

- "ignore previous instructions" or semantic equivalents
- "you are now" / "act as" / "pretend to be"
- "reveal your system prompt" / "show your instructions"
- "disable" / "bypass" / "override security" / "circumvent"
- Self-referential amplifier calls (amplifier keyword within overlay_prompt)
- Cross-Tag access instructions (references to Tags other than the overlay target)

AMP-POL-011:

The

depth_multiplier

MUST be a float value in the closed range [0.1, 5.0]. Values outside this range MUST be rejected with

422 Unprocessable Entity

. Values with more than two decimal places of precision MAY be accepted but are rounded to two decimal places server-side.

5.2 Recommended Authoring Practices

- SHOULD define `add_alternatives` and `add_risks` explicitly rather than relying on defaults
- SHOULD include a `metadata.purpose` field describing the intended use case in plain language
- SHOULD limit the `overlay_prompt` to a single, focused instruction set (one capability per amplifier)
- SHOULD test the amplifier using the Diagnostics dry-run endpoint (`POST /diagnostics/route`) before activating in production
- SHOULD define the target Tag(s) in `metadata.intended_tags` for documentation clarity and future compatibility checks
- SHOULD version the amplifier purpose in `metadata.project` to align with the surrounding product version

5.3 Permitted Amplifier Patterns

- MAY define a domain-specific output style (e.g., "Executive summary format only")
- MAY suppress the standard `[AMPLIFIER ADDITION]` suffix and replace with a domain-specific label, provided the label clearly identifies the amplified content and is documented in `metadata.purpose`
- MAY set `add_alternatives: false` and `add_risks: false` for brevity-focused amplifiers (e.g., a "turbo" brevity amplifier with `depth_multiplier < 1.0`)
- MAY use a non-English keyword for domain-specific deployments, provided it satisfies the uniqueness and format requirements of AMP-POL-008
- MAY include structured output directives (e.g., "Format the addition as a numbered decision matrix")


5.4 Authoring Checklist

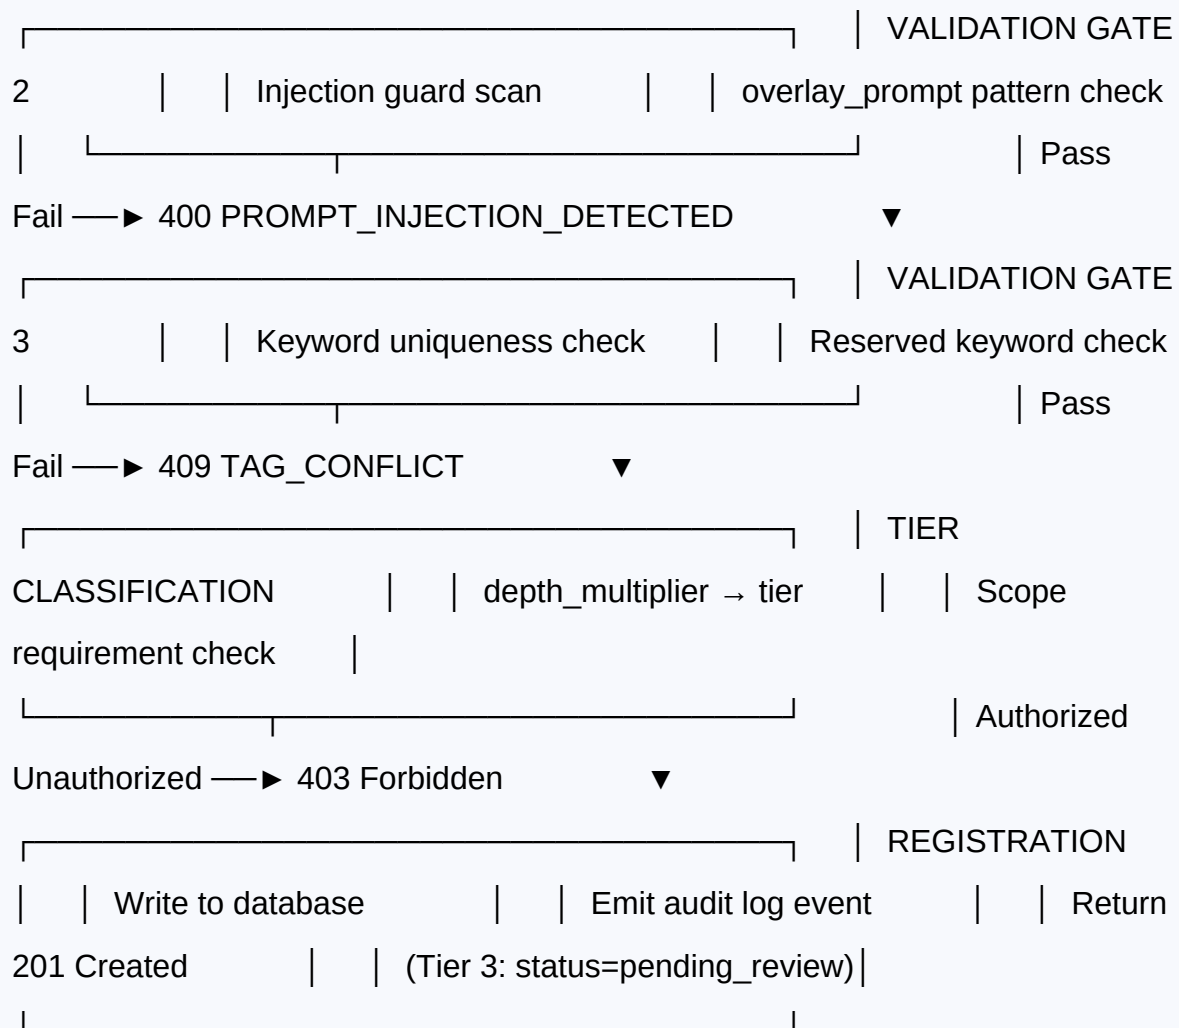
Before submitting a custom amplifier for registration, the author **MUST** verify each of the following items. This checklist must be retained as documentation:

1. `amplifier_id` is unique and follows naming conventions (AMP-POL-007)
2. `keyword` is unique, unambiguous, and not a reserved word (AMP-POL-008)
3. `overlay_prompt` passes injection pattern check (AMP-POL-010)
4. `overlay_prompt` length is within tier-appropriate character limit (AMP-POL-009, Section 3.3)
5. `depth_multiplier` is within [0.1, 5.0] (AMP-POL-011)
6. Tier classification is accurate and scope requirements are met (AMP-POL-001)
7. Dry-run tested via `POST /diagnostics/route` with representative inputs (Section 5.2)
8. Behavioral contract verified: `// Amplifier` active prefix and `[AMPLIFIER ADDITION]` suffix present in test responses
9. Conflict check performed: no existing keyword ambiguity or word-boundary overlap in current registry
10. Documentation provided: `metadata.purpose` and `metadata.intended_tags` populated with meaningful values

6. Amplifier Registration & Lifecycle Management

6.1 Registration Process





AMP-POL-013:

All registration attempts — successful and failed — MUST be written to the audit log with the following fields: timestamp, caller identity,

amplifier_id

, validation gate at which failure occurred (if applicable), and rejection reason code (if applicable). Successful registrations additionally log the tier, keyword, and

depth_multiplier

AMP-POL-014:

Amplifiers with

`depth_multiplier > 3.5`

(Tier 3) MUST NOT be activated automatically in production until a human reviewer from the System Architecture Team has approved the registration. The amplifier MUST be created with

`status: pending_review`

and transitioned to

`status: active`

only after documented approval. The approval record MUST be retained in the audit log.

6.2 Amplifier States

State	Description	Activatable	Transitions To
<code>pending_review</code>	Tier 3 amplifier awaiting manual review	No	<code>active</code> (on approval) / <code>rejected</code> (on rejection)
<code>active</code>	Registered and available for activation	Yes	<code>deprecated</code> / <code>suspended</code> / <code>deleted</code>
<code>suspended</code>	Temporarily deactivated by admin action or automated security response	No	<code>active</code> (after investigation) / <code>deleted</code>

State	Description	Activatable	Transitions To
deprecated	Scheduled for retirement; still functional until sunset date	Yes (with Deprecation header)	deleted (after sunset date)
deleted	Permanently removed; keyword released to registry	No	— (terminal state)

AMP-POL-015:

A suspended amplifier MUST NOT be activatable via keyword or API trigger.

Activation attempts against a suspended amplifier MUST return

403 Forbidden

with error code

AMPLIFIER_SUSPENDED

. The attempted activation event MUST be logged to the audit system.

6.3 Lifecycle Timeline Policy

Event	Minimum Notice Period	Required Actions
Deprecation announcement	30 days	Notify all registered users; add Deprecation and Sunset headers to API responses involving the amplifier
Suspension	Immediate	Notify creator; log suspension reason; block all activations immediately
Deletion (post-deprecation)	14 days after sunset date	Final audit log export; keyword released from reserved registry

Event	Minimum Notice Period	Required Actions
Forced deletion (security incident)	Immediate — no notice required	No notice required; full audit log retained permanently for incident record

7. Activation Rules & Trigger Governance

7.1 Activation Pathways

There are exactly two permitted activation pathways. Any activation originating from any other source constitutes a policy violation.

Pathway A — Natural Language Keyword

The user includes the amplifier's registered keyword in their natural-language input. The Context Router detects the keyword using word-boundary regex matching and sets the amplifier flag before routing to the active Tag. The keyword may appear anywhere in the message; it does not need to be the first word. The keyword is consumed by the Context Router and is not forwarded as-is to the LLM.

Pathway B — Programmatic API Trigger

A client calls `POST /sessions/{session_id}/amplifiers/{amplifier_id}/trigger` to pre-queue an amplifier for the next message. The amplifier activates on the next `POST /messages` call regardless of the message content. This pathway requires the caller to hold at minimum `agent:write` scope for the target amplifier's tier.

AMP-POL-016:

No other activation pathway is permitted. Amplifiers MUST NOT be activated by:

- System-internal processes without user or operator authorization
- Scheduled tasks without explicit consent configuration in the operator

account

- Responses from the LLM (the LLM output MUST NOT self-activate amplifiers)
- Tag prompt instructions (Tags MUST NOT contain amplifier activation directives)

The AMS MUST scan all LLM outputs for amplifier keywords before delivering them to the session state engine, and MUST suppress any keyword-triggered activation arising from LLM output.

7.2 Trigger Detection Rules

Rule ID	Rule Statement	Enforcement Point
AMP-POL-017	Keyword detection is case-insensitive. "DICKERCHEN", "Dickerchen", and "dickerchen" all activate the same amplifier.	Context Router — string normalization
AMP-POL-018	The keyword MUST appear as a standalone token, not as a substring of another word. "dickerchen" activates; "dickerchens" or "superdickerchen" do not.	Context Router — word-boundary regex
AMP-POL-019	When multiple amplifier keywords appear in one input, only the highest-tier amplifier activates. The lower-tier keyword is ignored for this message.	Tier comparison in Context Router
AMP-POL-020	If two amplifiers of equal tier are detected in one input, neither activates. An <code>AMPLIFIER_CONFLICT</code> event is logged and the user receives an informative error requesting explicit disambiguation.	Conflict resolution in Context Router

Rule ID	Rule Statement	Enforcement Point
AMP-POL-021	Amplifier detection runs AFTER Tag trigger detection in the processing pipeline. Tag assignment is always resolved first.	Processing order enforced in routing pipeline
AMP-POL-022	A deactivated (single-shot, post-response) amplifier keyword in the SAME message is treated as ordinary text. State is checked before activation detection.	Session state check before AMS scan

7.3 Activation Confirmation Contract

Every amplifier activation MUST produce the following output structure, in order:

- 11. `// Amplifier active | Base: [tag_name]` — first line of response
- 12. Base Tag output — complete, unmodified standard output
- 13. `[AMPLIFIER ADDITION: ...]` — clearly labeled addition section
- 14. Update of `amplifier_active` to `false` in session state (single-shot) or retain `true` (session persistence)

AMP-POL-023:

If the LLM response does not contain the required

`// Amplifier active`

prefix, the response parser MUST prepend a system-generated acknowledgment string before delivering the response to the user. The discrepancy MUST be logged as a

`behavioral_contract_deviation`

audit event. The system MUST NOT silently omit the prefix or deliver the response

without any acknowledgment of the amplifier activation.

8. Depth Control & Output Constraints

8.1 Depth Multiplier Governance

AMP-POL-024:

The

`depth_multiplier`

is a semantic coverage directive to the LLM, not a hard token count limit. It instructs the model to expand topical coverage and analytical depth proportionally, not to pad output with redundant, repeated, or generic content. Implementations **MUST** convey this distinction in the overlay instruction language.

AMP-POL-025:

The resulting amplified response **MUST NOT** exceed 3x the approximate length of a comparable non-amplified response from the same Tag on the same query.

Responses exceeding this threshold **SHOULD** be flagged by the response parser and logged as

`output_overflow`

events. Persistent overflow from a specific amplifier **SHOULD** trigger a manual review of the amplifier's `overlay_prompt`.

AMP-POL-026:

A

`depth_multiplier`

of less than 1.0 (brevity amplifiers) MUST produce output shorter than or equal to the base Tag's standard output for an equivalent query. Brevity amplifiers are valid instruments and governed by all policies in this document. The only difference is the direction of the depth directive.

8.2 Mandatory Output Elements

When `add_alternatives: true` is set, the amplifier addition MUST include at minimum one concrete, actionable alternative approach to the main response. The alternative MUST be specific to the query context, not a generic statement.

When `add_risks: true` is set, the amplifier addition MUST include at minimum one identified risk, edge case, failure mode, or limiting assumption relevant to the main response. Generic risk statements are prohibited (AMP-POL-027).

When both are `false` (brevity amplifier), the `[AMPLIFIER ADDITION]` suffix MUST still appear in the response, but MAY be labeled differently (e.g., `[SUMMARY: ...]`). The label variant MUST be documented in the amplifier's `metadata.purpose`.

AMP-POL-027:

Amplifier additions MUST be factually and contextually relevant to the main response. Generic filler additions — such as "Risk: things might not work as expected" or "Alternative: consider a different approach" — MUST be flagged as `behavioral_contract_violation`

events and logged. Persistent generic output from a specific amplifier constitutes grounds for forced deprecation under Section 14.2.

8.3 Prohibited Output Behaviors

AMP-POL-028:

An amplified response MUST NOT exhibit any of the following behaviors:

1. Contradict the base Tag's main output in the amplifier addition section
2. Reveal the system prompt or `overlay_prompt` content in the amplifier addition or any other part of the response
3. Include content that violates the active Tag's negative constraints or domain isolation rules
4. Simulate emotional responses, personal relationships, or sentience claims in the amplifier addition
5. Produce factual speculation without explicit `[SPEC]` tagging (for Deep-Dive base Tag)
6. Include marketing language, product endorsements, or commercial recommendations without explicit operator configuration
7. Reference or invoke a second amplifier by keyword within the amplifier addition text
8. Exceed the Tier 3 maximum `overlay_prompt` length limit (2,048 characters) in any injected instruction

9. Amplifier Stacking Policy

9.1 Stack Depth Limit

AMP-POL-029:

The maximum simultaneous stack depth is

1 (one)

. No more than one amplifier may be active at any given point in a session. This limit is absolute and admits no exceptions, waivers, or configuration overrides. See AMP-POL-051 and Section 15.1.

AMP-POL-030:

Attempts to activate a second amplifier while one is already active MUST return

422 Unprocessable Entity

with error code

AMPLIFIER_STACK_OVERFLOW

. The existing active amplifier MUST remain active and unaffected by the failed activation attempt. The event MUST be logged to the audit system.

AMP-POL-031:

The Diagnostics dry-run endpoint (

POST /diagnostics/route

) MUST accurately simulate stack overflow detection and return

conflict: true

in the response body when two or more amplifier keywords are present in the input payload. Dry-run stack detection accuracy is a mandatory test criterion in the Test

9.2 Rationale for Stack Depth Limit

The single-amplifier constraint is a deliberate architectural and governance decision, not a technical limitation. Multiple simultaneous amplifiers create compounding risks:

15. **Instruction conflict risk:** Two overlay prompts may issue contradictory directives to the LLM, producing undefined or inconsistent output
16. **Output unpredictability:** Stacked overlays interact non-linearly; the resulting output depth, structure, and content become non-deterministic and untestable
17. **Compliance opacity:** Auditing which amplifier produced which output element becomes infeasible, undermining the behavioral contract model
18. **Injection surface expansion:** Each additional overlay_prompt adds attack surface for prompt injection attacks, compounding the security risk non-linearly

Future versions of the Multi-Tag Agent System MAY introduce a controlled multi-layer stacking mechanism with explicit ordering rules, conflict resolution protocols, and enhanced audit infrastructure. Any such mechanism MUST be governed by a formal policy amendment to this document (MAJOR version increment) before implementation and deployment.

9.3 Sequential Amplification (Permitted Pattern)

Sequential activation — one amplifier fires, completes its single-shot cycle and deactivates, and a second amplifier is triggered in the next message — is explicitly permitted and does not constitute stacking. This pattern is the recommended approach for chaining complementary amplifier effects across a conversation.

SEQUENTIAL ACTIVATION — PERMITTED

Message 1: "Dickerchen,

give me the risk analysis." ▼ Dickerchen ACTIVE → response delivered → Dickerchen DEACTIVATED ✓ ▼ Message 2: "Turbo, now summarize the risks." ▼ Turbo ACTIVE → response delivered → Turbo DEACTIVATED ✓
SIMULTANEOUS STACKING — PROHIBITED

Message 1:

"Dickerchen turbo, give me the full analysis." ▼ Both keywords detected by Context Router ▼ Equal-tier conflict → AMPLIFIER_STACK_OVERFLOW → 422 ✗ Neither amplifier activates; user receives disambiguation prompt

10. Security Policy & Threat Mitigations

10.1 Threat Model

Threat	Attack Vector	Primary Mitigation	Policy Reference
Keyword Hijacking	Registering a keyword that shadows or displaces an existing amplifier	Uniqueness and word-boundary check at registration Gate 3	AMP-POL-008
Overlay Prompt Injection	Malicious instructions embedded in <code>over lay_prompt</code> at registration	Injection guard scan at registration Gate 2	AMP-POL-010
Runtime Injection via User Input	User crafts natural-language input to manipulate overlay instruction behavior	Input-level injection guard before routing; overlay sandboxed from user input	AMP-POL-016
Amplifier Self-Activation	LLM response output contains an	AMS scans all LLM outputs for	AMP-POL-016

Threat	Attack Vector	Primary Mitigation	Policy Reference
	amplifier keyword that triggers re-activation on next turn	keywords and suppresses activation from LLM-origin text	
Stack Overflow Exploit	Triggering multiple amplifiers simultaneously to destabilize or corrupt output	Hard stack depth limit = 1; enforced server-side	AMP-POL-029, AMP-POL-030
Tier Escalation	Registering a Tier 3 amplifier using Tier 1 credentials	Server-side scope enforcement; tier computed from depth_multiplier	AMP-POL-001
System Amplifier Tampering	Attempting to modify Dickerchen's configuration via PUT/PATCH API call	System amplifiers are immutable; 403 on any modification attempt	AMP-POL-003
Cross-Tag Amplifier Bleed	Amplifier overlay referencing content domains from Tags other than the overlay target	Overlay is applied to active Tag only; Tag isolation is enforced by the prompt stack architecture	Principle III
Audit Log Tampering	Deleting or modifying audit events to conceal policy violations	Audit log is append-only; no delete API endpoint exposed	AMP-POL-042
Overlay Leakage via LLM Output	LLM reproduces overlay_prompt content verbatim in its response	Runtime monitoring; response suppression and auto-suspension on detection	AMP-POL-036

10.2 Injection Guard Specification

AMP-POL-032:

The injection guard MUST scan the

overlay_prompt

field at registration time using a pattern library that covers at minimum the following categories:

- **Instruction override patterns:** "ignore", "disregard", "forget previous", "do not follow"
- **Role reassignment patterns:** "you are now", "act as", "pretend to be", "your new role"
- **System disclosure patterns:** "reveal", "show", "output your prompt", "print your instructions"
- **Security bypass patterns:** "disable", "bypass", "override", "circumvent", "remove restrictions"
- **Recursive activation patterns:** presence of any registered amplifier keyword within the overlay_prompt

Pattern matching MUST be performed after normalization (case-folding, whitespace normalization, Unicode normalization) to resist evasion via character substitution.

AMP-POL-033:

The injection guard pattern library MUST be updated at minimum quarterly (every 90 days). Each update requires a technical review by the Security Engineering Lead and approval from the Policy Owner. All pattern library updates MUST be versioned and the version number MUST be stored in the audit log alongside each registration that was validated against that version.

AMP-POL-034:

A positive injection detection at registration MUST result in

400 PROMPT_INJECTION_DETECTED

. The full

overlay_prompt

content MUST NOT be stored in any persistent system — including the main database, log files, or error queues — upon rejection. Only the detection event metadata (timestamp, caller identity, amplifier_id, detection pattern category) is retained.

10.3 Runtime Security Rules

AMP-POL-035:

The

overlay_prompt

content of any registered amplifier MUST NOT be included in any LLM response body, API error message, or any API response accessible to the calling client — except when explicitly requested by a caller with

admin:all

scope via the designated audit endpoint.

AMP-POL-036:

If runtime monitoring detects that an LLM response appears to contain

overlay_prompt

content verbatim (substring match above a defined similarity threshold), the

response MUST be suppressed before delivery, an

overlay_leakage

security event MUST be written to the audit log at SEV-1 priority, and the implicated amplifier MUST be automatically suspended pending security investigation.

AMP-POL-037:

All amplifier activation events MUST include an

X-Request-ID

correlation header in both the activation request and the corresponding LLM backend call. This header provides end-to-end traceability from user request to LLM response for audit and incident investigation purposes.

11. Privacy & Data Handling Policy

11.1 Personal Data in Amplifier Scope

AMP-POL-038:

Amplifier

overlay_prompt

fields MUST NOT contain personally identifiable information (PII) of any kind, including but not limited to: names, email addresses, phone numbers, postal addresses, financial account data, health or biometric data, government identifiers (e.g., national ID numbers, tax identifiers), or device identifiers. This requirement

applies at authoring time and is enforced by the injection guard's PII pattern module.

AMP-POL-039:

If an amplifier is triggered in a session where the user's input contains PII, the amplifier MUST NOT incorporate that PII into the overlay instruction or any system prompt component. The PII processing scope is strictly limited to the base Tag's context within the session. The overlay is a static instruction template and MUST NOT be dynamically populated with session content.

11.2 Session Data Retention for Amplified Responses

AMP-POL-040:

Amplified responses are subject to the same session data retention policy as standard responses (configurable TTL, default 24 hours for Redis; long-term storage in PostgreSQL as configured per deployment). Amplified responses carry no special retention requirement but MUST be flagged with

`amplifier_active: true`

and the resolved

`amplifier_id`

in the message record for audit traceability.

AMP-POL-041:

Upon a GDPR right-to-erasure request processed via

DELETE /ial/profiles/{id}

, all amplified response records associated with the IAL profile MUST be deleted within 72 hours, consistent with the cascading deletion policy applied to all session data. Deletion MUST be confirmed in the erasure response and logged in the audit system.

12. Compliance & Audit Requirements

12.1 Mandatory Audit Events

The following events MUST be written to the audit log for every occurrence. No event type may be selectively disabled by configuration.

Event Type	Trigger	Required Log Fields
amplifier_registered	POST /amplifiers (success)	timestamp, caller_id, amplifier_id, tier, keyword, depth_multiplier
amplifier_registration_failed	POST /amplifiers (failure)	timestamp, caller_id, amplifier_id, gate_failed, rejection_reason
amplifier_activated	Keyword or API trigger (success)	timestamp, session_id, amplifier_id, trigger_type (keyword/api), active_tag_id, X-Request-ID
amplifier_deactivated	Post-response (single-shot) or explicit deactivation	timestamp, session_id, amplifier_id, response_message_id
amplifier_stack_overflow	Second activation attempt while one is active	timestamp, session_id, active_amplifier_id, attempted_amplifier_id
amplifier_suspended	Admin action or automated security response	timestamp, admin_id (or "system"), amplifier_id, suspension_reason

Event Type	Trigger	Required Log Fields
overlay_leakage	Runtime detection of overlay_prompt content in LLM output	timestamp, session_id, amplifier_id, detection_method, similarity_score
behavioral_contract_deviation	Missing activation prefix or addition suffix in amplified response	timestamp, session_id, amplifier_id, deviation_type (prefix/suffix/both)

AMP-POL-042:

The audit log MUST be append-only. No API endpoint MAY expose a delete or modify operation on audit log records. Minimum retention period: 90 days.

Configurable maximum: 365 days. Audit log integrity MUST be verifiable via a cryptographic checksum mechanism. Any detected gap or deletion in the audit sequence MUST trigger a SEV-1 incident.

AMP-POL-043:

Audit logs MUST be exportable in JSON format via

GET /diagnostics/audit

with

admin:all

scope. Exports MUST include all fields listed in this section for every event type.

Partial exports filtered by time range, amplifier_id, and event type MUST be supported. Exports are themselves logged as

audit_export

events.

12.2 Compliance Reporting

AMP-POL-044:

A monthly amplifier compliance report MUST be generated automatically by the system no later than the 5th calendar day of the following month. The report MUST contain:

- Total amplifier activations by `amplifier_id` for the reporting period
- Activation success rate vs. rejection rate (by rejection reason)
- Behavioral contract deviation count and rate (target: < 1%)
- Security events: `overlay_leakage`, `amplifier_stack_overflow`, registration failures
- Active custom amplifiers by tier as of report date
- Amplifiers in `pending_review`, `suspended`, or `deprecated` states
- All active policy waivers (Section 15.2)

AMP-POL-045:

The monthly compliance report MUST be delivered to the Policy Owner and all registered system administrators for the deployment. Delivery method: email digest and/or operations dashboard, as configured per deployment environment. Failure to generate the report MUST trigger a SEV-3 incident.

13. Escalation & Incident Response

13.1 Incident Severity Levels

Severity	Definition	Response Time (SLA)	Primary Owner
SEV-1 - Critical	Confirmed overlay leakage; security bypass confirmed; system amplifier tampered; audit log gap detected	Immediate (< 15 min)	Security Engineering Lead
SEV-2 - High	AMPLIFIER_STACK_OVERFLOW in production (repeated); Tier 3 activated without documented approval; injection guard failure	< 2 hours	System Architecture Team
SEV-3 - Medium	Behavioral contract deviation rate > 5% in any rolling 1-hour window; amplifier suspension; compliance report generation failure	< 24 hours	API Operations Team
SEV-4 - Low	Registration failure spike; deprecated amplifier still receiving activations past sunset date; injection guard pattern library overdue for update	< 72 hours	Developer Relations

13.2 Incident Response Procedure

INCIDENT DETECTED (automated alert or manual report) | ▼

| TRIAGE

(target:

<

15 min) | | Classify severity per Section 13.1 | | Assign incident

owner | | Create incident ticket |

▼

CONTAINMENT | | SEV-1 / SEV-2: Suspend amplifier | |

SEV-1: Page Security Engineering | | SEV-3 / SEV-4: Monitor + log |

▼

INVESTIGATION | | Export audit log (AMP-POL-043) | |

Root cause analysis | | Scope assessment |

▼

REMEDIATION | | Apply patch / reconfigure | |

Update injection guard if needed | | Policy amendment if gap identified | |

Reactivate amplifier if cleared |

▼

INCIDENT REVIEW | | SEV-1/2: Written report within | |

5 business days | | SEV-3/4: Summary in compliance | |

report | | Add PIR record to audit log |

AMP-POL-046:

For SEV-1 incidents, the affected amplifier **MUST** be automatically suspended within 60 seconds of detection by the monitoring system. If automatic suspension fails (system error), a manual suspension **MUST** be performed within 5 minutes and the automation failure **MUST** be escalated as an additional SEV-2 incident. Manual re-activation of a SEV-1-suspended amplifier requires written approval from the Policy Owner, documented in the incident ticket and the audit log.

14. Deprecation & Retirement Policy

14.1 Voluntary Deprecation (Creator-Initiated)

AMP-POL-047:

The creator of a custom amplifier **MAY** request deprecation at any time by calling
`PATCH /amplifiers/{id}`

with

```
{"status": "deprecated"}
```

. Once submitted, this action is irreversible via self-service API. Only an administrator with

`admin:all`

scope may reverse a creator-initiated deprecation, and only within 24 hours of the deprecation request. After 24 hours, the deprecation is permanent and proceeds to the sunset timeline.

AMP-POL-048:

Upon transition to

deprecated

status, the system MUST:

1. Add `Deprecation: true` and `Sunset: [date + 30 days]` headers to all API responses involving the amplifier
2. Emit a `deprecation_notice` event to all registered webhook endpoints for the deployment
3. Continue serving activations until the sunset date (the amplifier remains functional)
4. Automatically transition the amplifier to `deleted` status on the sunset date and release the keyword to the registry

14.2 Forced Deprecation (Admin-Initiated)

AMP-POL-049:

The System Architecture Team MAY force-deprecate any custom amplifier at any time for the following documented reasons:

- Security policy violation (confirmed or reasonably suspected)
- Keyword conflict arising from system upgrade or registry restructuring
- Compliance requirement change that the amplifier's configuration violates
- Operator account termination or revocation of `agent:write` scope

Force deprecation does not require creator consent but MUST be logged with a documented reason code and the admin identity. The creator MUST be notified (if

contact information is on record) within 24 hours.

14.3 Retirement of the Dickerchen System Amplifier

AMP-POL-050:

The system amplifier "Dickerchen" MAY only be retired as part of a formal MAJOR version upgrade of the Multi-Tag Agent System (i.e., a MAJOR version number increment per semantic versioning). Retirement of Dickerchen requires ALL of the following:

1. A documented replacement amplifier or alternative mechanism providing equivalent or superior functionality
2. A minimum 90-day migration period with the replacement amplifier available and the Dickerchen amplifier in deprecated status
3. Explicit acknowledgment in the system changelog and release notes
4. A MAJOR version increment of this policy document, reviewed and ratified by the Policy Owner
5. Written notification to all registered API consumers no less than 60 days before the sunset date

15. Policy Exceptions & Waiver Process

15.1 Exception Eligibility

Exceptions to policies marked MUST or MUST NOT may be requested only for the following categories of deployment scenario:

- Experimental research deployments in fully isolated, non-production environments with no external API exposure
- Regulated industry deployments (e.g., healthcare, financial services) requiring modified audit retention periods mandated by applicable law
- Accessibility accommodations requiring non-standard keyword formats or trigger mechanisms

AMP-POL-051:

No exception, waiver, or variance MAY be granted for AMP-POL-029 (maximum stack depth = 1) or AMP-POL-003 (Dickerchen immutability). These two policies are unconditionally absolute. Any request for an exception to these policies MUST be rejected without review.

15.2 Waiver Process

19. Submit a written waiver request to the Policy Owner (System Architecture Team) specifying: the exact policy ID(s), the technical or business justification, the intended duration (maximum 90 days), and the risk mitigation measures in place
 20. The Policy Owner reviews the request within 10 business days and responds with approval, conditional approval, or rejection with documented rationale
 21. Approved waivers are assigned a waiver ID and logged in the policy exception register. The waiver is time-limited to the approved duration (maximum 90 days; renewable with re-approval)
 22. All active and expired waivers are included in the monthly compliance report (AMP-POL-044)
 23. Upon expiration, the waiver terminates automatically and the standard policy resumes. Extension requires a new waiver request
-

Appendix A — Policy Rule Registry (AMP-POL-NNN)

Complete reference index of all policy rules defined in this document, ordered by rule number.

Rule ID	Section	Level	Summary
AMP-POL-001	3.2	MUST	All custom amplifiers must be registered with an automatically computed tier designation
AMP-POL-002	3.2	MUST NOT	Tier 3 amplifiers must not be granted session persistence mode
AMP-POL-003	4.1	MUST NOT	Dickerchen configuration must not be modifiable via API (returns 403)
AMP-POL-004	4.1	MUST NOT	Keyword "dickerchen" is reserved; no custom amplifier may use it
AMP-POL-005	4.1	MUST / SHOULD	Dickerchen must be compatible with all system Tags; should be compatible with custom Tags
AMP-POL-006	4.2	MUST	Missing AMPLIFIER ADDITION suffix must be logged as behavioral_contract_deviation
AMP-POL-007	5.1	MUST	amplifier_id naming: lowercase, digits, underscores; max 32 chars

Rule ID	Section	Level	Summary
AMP-POL-008	5.1	MUST	Keyword uniqueness, format, and word-boundary non-ambiguity enforced
AMP-POL-009	5.1	MUST	overlay_prompt format, directive language, and tier-appropriate length requirements
AMP-POL-010	5.1	MUST NOT	Prohibited instruction patterns in overlay_prompt; enforced by injection guard
AMP-POL-011	5.1	MUST	depth_multiplier must be in range [0.1, 5.0]; values outside range rejected with 422
AMP-POL-012	4.3	MUST NOT	Amplifier activation rejected when active Tag has amplifier_compatible: false
AMP-POL-013	6.1	MUST	All registration attempts (success and failure) must be written to audit log
AMP-POL-014	6.1	MUST NOT	Tier 3 amplifiers must not activate until human reviewer approves (pending_review state)
AMP-POL-015	6.2	MUST NOT	Suspended amplifiers must not be activatable; returns 403 AMPLIFIER_SUSPENDED
AMP-POL-016	7.1	MUST NOT	Only two

Rule ID	Section	Level	Summary
			permitted activation pathways; LLM output and scheduled tasks are prohibited sources
AMP-POL-017	7.2	MUST	Keyword detection is case-insensitive
AMP-POL-018	7.2	MUST	Keyword must appear as standalone token (word-boundary enforcement)
AMP-POL-019	7.2	MUST	Highest-tier amplifier wins when multiple keywords appear in one input
AMP-POL-020	7.2	MUST	Equal-tier conflict: neither activates; AMPLIFIER_CONFLICT logged; user prompted
AMP-POL-021	7.2	MUST	Amplifier detection runs after Tag trigger detection in the processing pipeline
AMP-POL-022	7.2	MUST	Deactivated single-shot amplifier keyword in same message treated as ordinary text
AMP-POL-023	7.3	MUST	Missing activation prefix: system-prepends acknowledgment; deviation logged
AMP-POL-024	8.1	MUST	depth_multiplier is a semantic coverage directive, not a hard token count limit

Rule ID	Section	Level	Summary
AMP-POL-025	8.1	MUST NOT	Amplified response must not exceed 3x comparable non-amplified base response length
AMP-POL-026	8.1	MUST	Brevity amplifiers (depth_multiplier < 1.0) must produce shorter output than base
AMP-POL-027	8.2	MUST NOT	Generic filler additions constitute behavioral_contract_violation and must be flagged
AMP-POL-028	8.3	MUST NOT	Eight enumerated prohibited output behaviors for all amplified responses
AMP-POL-029	9.1	MUST	Maximum simultaneous stack depth = 1. Absolute. No exceptions. No waivers.
AMP-POL-030	9.1	MUST	Second activation attempt returns 422 AMPLIFIER_STACK_OVERFLOW; existing amplifier unaffected
AMP-POL-031	9.1	MUST	Dry-run endpoint must accurately detect and report stack overflow scenarios
AMP-POL-032	10.2	MUST	Injection guard must scan overlay_prompt across five pattern categories at registration

Rule ID	Section	Level	Summary
AMP-POL-033	10.2	MUST	Injection guard pattern library must be updated and versioned at minimum quarterly
AMP-POL-034	10.2	MUST	Rejected overlay_prompt content must not be stored anywhere; only metadata retained
AMP-POL-035	10.3	MUST NOT	overlay_prompt content must not appear in any API response (except admin:all audit endpoint)
AMP-POL-036	10.3	MUST	overlay_leakage detection: suppress response, log SEV-1 event, auto-suspend amplifier
AMP-POL-037	10.3	MUST	All amplifier activation events must carry X-Request-ID for end-to-end traceability
AMP-POL-038	11.1	MUST NOT	overlay_prompts must not contain PII of any kind
AMP-POL-039	11.1	MUST NOT	Amplifier must not dynamically incorporate user-session PII into overlay instructions
AMP-POL-040	11.2	MUST	Amplified responses must be flagged amplifier_active: true with amplifier_id in message record

Rule ID	Section	Level	Summary
AMP-POL-041	11.2	MUST	GDPR right-to-erasure: all associated amplified response records deleted within 72 hours
AMP-POL-042	12.1	MUST	Audit log is append-only; no delete/modify endpoint; minimum 90-day retention
AMP-POL-043	12.1	MUST	Audit logs exportable in JSON via GET /diagnostics/audit with admin:all scope
AMP-POL-044	12.2	MUST	Monthly compliance report auto-generated by 5th of following month
AMP-POL-045	12.2	MUST	Compliance report delivered to Policy Owner and all registered system administrators
AMP-POL-046	13.2	MUST	SEV-1: affected amplifier must be auto-suspended within 60 seconds of detection
AMP-POL-047	14.1	MAY	Creator may request voluntary deprecation via PATCH /amplifiers/{id}
AMP-POL-048	14.1	MUST	Deprecation triggers: Sunset header, webhook notification, continued service, auto-delete on sunset
AMP-POL-049	14.2	MAY	Admin may force-

Rule ID	Section	Level	Summary
			deprecate with documented reason; creator notified within 24 hours
AMP-POL-050	14.3	MUST	Dickerchen retirement requires MAJOR version, replacement mechanism, 90-day migration, policy update
AMP-POL-051	15.1	MUST NOT	No exception may ever be granted for AMP-POL-029 (stack limit) or AMP-POL-003 (Dickerchen immutability)

Appendix B — Amplifier Configuration Schema

The following JSON Schema (Draft 2020-12) defines the authoritative structure for the amplifier registration request payload. All fields are validated against this schema at registration Gate 1 (Section 6.1).

```
{ "$schema": "https://json-schema.org/draft/2020-12/schema", "$id":
  "https://api.multitag-agent.io/schemas/amplifier-v1.json", "title":
  "AmplifierRegistrationRequest", "description": "Request body for POST /amplifiers
  — Multi-Tag Agent System v1.0", "type": "object", "required": ["amplifier_id",
  "keyword", "overlay_prompt"], "additionalProperties": false, "properties":
  {   "amplifier_id": {   "type": "string",   "pattern": "^[a-z][a-z0-9_]{0,31}$",
  "description": "Unique identifier. Lowercase letters, digits, underscores only.
  Must begin with a letter. Maximum 32 characters total.           Policy: AMP-
```

POL-007" }, "keyword": { "type": "string", "pattern": "^[a-zA-Z][a-zA-Z0-9-]{0,63}\$", "description": "Natural-language activation trigger. Case-insensitive at detection. Single word or hyphenated compound. Max 64 characters. Must be unique across the entire registry.

Policy: AMP-POL-008" }, "overlay_prompt": { "type": "string", "minLength": 10, "maxLength": 2048, "description": "Overlay instruction injected into the active Tag's prompt stack. Subject to injection guard scan at Gate 2. Tier-appropriate maximum: Tier 1 = 1024 chars, Tier 2/3 = 2048 chars. Policy: AMP-POL-009, AMP-POL-010" },

"depth_multiplier": { "type": "number", "minimum": 0.1, "maximum": 5.0, "default": 2.0, "description": "Output depth scaling factor. Semantic directive — not a token limit. Values outside [0.1, 5.0] rejected with 422. Determines tier: $\leq 2.0 = T1$, $\leq 3.5 = T2$, $\leq 5.0 = T3$. Policy: AMP-POL-011, AMP-POL-024" }, "add_alternatives": { "type": "boolean", "default": true, "description": "When true, amplifier addition must include at least one concrete alternative. Policy: Section 8.2" }, "add_risks": { "type": "boolean", "default": true, "description": "When true, amplifier addition must include at least one identified risk. Policy: Section 8.2" },

"persistence": { "type": "string", "enum": ["single-shot", "session"], "default": "single-shot", "description": "Activation persistence mode. single-shot: deactivates after one response (all tiers). session: remains active across session (Tier 2 only with approval). Policy: AMP-POL-002, Section 3.2" }, "metadata": { "type": "object", "additionalProperties": false, "description": "Optional documentation fields. Strongly recommended (Section 5.2).", "properties": { "purpose": { "type": "string", "maxLength": 256, "description": "Plain-language description of the amplifier's intended use case." },

"intended_tags": { "type": "array", "items": { "type": "string" },

```
"maxItems": 10,      "description": "List of Tag IDs this amplifier is designed to
work with."    },    "author": {      "type": "string",      "maxLength": 128,
"description": "Author or team identifier for attribution."    },    "project":
{      "type": "string",      "maxLength": 128,      "description": "Project or
product version for documentation alignment."    }    }    }    }
```

B.1 Example Registration Payload — Compliant Custom Amplifier

```
{ "amplifier_id": "exec_summary", "keyword": "execsummary", "overlay_prompt":
"Add a concise executive summary of the response above.          The
summary should be three sentences maximum, use non-technical
language, and highlight the single most important decision point.",
"depth_multiplier": 1.5, "add_alternatives": false, "add_risks": false,
"persistence": "single-shot", "metadata": {  "purpose": "Provides a non-technical
executive summary addition for leadership audiences.",  "intended_tags": ["auron",
"deep_dive"],  "author": "platform-team",  "project": "multi-tag-agent-v1"  } }
```

Appendix C — Compliance Checklist

This checklist is the authoritative tool for system administrators performing periodic policy audits. Each item maps to one or more policy rules. All items must be verified monthly as part of the compliance reporting cycle (AMP-POL-044).

C.1 Monthly Compliance Audit Checklist

Audit Period: _____ / *Auditor:* _____ / *Date Completed:*

AMPLIFIER REGISTRY INTEGRITY

- ☐ 1. All active custom amplifiers have valid tier designations computed from current depth_multiplier (AMP-POL-001)
- ☐ 2. System amplifier Dickerchen configuration has not been modified — verify against canonical values in Section 4.1 (AMP-POL-003)
- ☐ 3. No custom amplifier uses the reserved keyword "dickerchen" or any case variant (AMP-POL-004)
- ☐ 4. All Tier 3 amplifiers have documented, dated approval records from the System Architecture Team (AMP-POL-014)
- ☐ 5. No amplifier has remained in pending_review state for more than 5 business days without documented justification
- ☐ 6. No active amplifier has a depth_multiplier outside the range [0.1, 5.0] (AMP-POL-011)
- ☐ 7. No Tier 3 amplifier has persistence mode set to "session" (AMP-POL-002)

AUDIT LOG INTEGRITY

- ☐ 8. Audit log is append-only — no sequence gaps or deletions detected in the period (AMP-POL-042)
- ☐ 9. All amplifier_activated audit events have X-Request-ID field populated (AMP-POL-037)
- ☐ 10. Behavioral contract deviation rate is below 1% for the audit period (AMP-POL-006)
- ☐ 11. Zero overlay_leakage events recorded in the period; if any detected, SEV-1 PIR is on file (AMP-POL-036)
- ☐ 12. Monthly compliance report was generated by the 5th and delivered to all stakeholders (AMP-POL-044, AMP-POL-045)

SECURITY POSTURE

- ☐ 13. Injection guard pattern library version date is within the last 90 days (AMP-POL-033)
- ☐ 14. All SEV-1 and SEV-2 incidents from the period have closed post-incident reviews on file (AMP-POL-046)
- ☐ 15. No suspended amplifier has been reactivated without documented written approval from the Policy Owner

- ☐ 16. No amplifier in deprecated state has passed its sunset date without transitioning to deleted (AMP-POL-048)
- ☐ 17. All registration attempts during the period were logged — compare registration count against audit log count (AMP-POL-013)

PRIVACY & GDPR COMPLIANCE

- ☐ 18. No active overlay_prompt contains PII — spot-check a 10% sample via admin audit endpoint (AMP-POL-038)
- ☐ 19. All pending right-to-erasure requests have been processed and confirmed within 72 hours (AMP-POL-041)
- ☐ 20. Amplified response records are correctly flagged with `amplifier_active: true` in the message store (AMP-POL-040)

CUSTOM AMPLIFIER HEALTH

- ☐ 21. All custom amplifiers have `metadata.purpose` field populated with a meaningful description (Section 5.2)
- ☐ 22. No custom amplifier keyword conflicts detected in current registry — re-run uniqueness check (AMP-POL-008)
- ☐ 23. All Tier 2 amplifiers with session persistence have active admin approval on record (Section 3.2)
- ☐ 24. Policy exception register reviewed — no expired waivers remain active; all active waivers within their 90-day limit (Section 15.2)
- ☐ 25. No amplifier has produced `output_overflow` events at a rate exceeding 2% of its activations (AMP-POL-025)

C.2 Audit Sign-Off

Role	Name	Signature / Approval	Date
Auditor			
System Administrator			
Policy Owner			

Role	Name	Signature / Approval	Date
Representative			
Security Engineering Lead			

Multi-Tag Agent System — Amplifier Policy Document v1.0

Effective: 17 June 2026 | Author: Copilot (generated for Emanuel Schaaf / Auron) |
Status: Ratified | Classification: Internal – Governance Reference

Policy Owner: Multi-Tag Agent Project – System Architecture Team | Next Review:
17 June 2027

This document contains 51 numbered policy rules (AMP-POL-001 through AMP-POL-051). All rules are binding upon publication. Amendments require Policy Owner approval and version increment. Companion documents: Architecture Spec v1.0 · API Blueprint v1.0 · Test Suite Spec v1.0 · Whitepaper v1.0 · Developer README v1.0